

122022010 Shrayank Jai Mistry

Description: Implementation of ARIES Algorithm.

Programming Language: Python

```
1  from tabulate import tabulate
2  import copy
3
4  /* Keep track of the LSN no. for logs
5  LSN = 0
6
7  /* Storage for Log Records(On disk and on Memory)
8  LOG_DISK = {}
9  LOG_MEMORY = {}
10
11 /* Keep track of lastLSN flushed from memory to disk
12 flushedLSN = None
13
14 /* Keep track of the Last Completed checkpoint
15 masterRecord = [0, 0]
16
17 /* Class for log records
18 class LogRecord:
19     def __init__(self, LSN, prevLSN, TxnId, type, dataItem, before, after, undoNext):
20         self.LSN = LSN
21         self.prevLSN = prevLSN
22         self.TxnId = TxnId
23         self.type = type
24         self.dataItem = dataItem
25         self.before = before
26         self.after = after
27         self.undoNext = undoNext
28
29     def get_record(self):
30         return [str(self.LSN), str(self.prevLSN), str(self.TxnId), str(self.type), str
                (self.dataItem), str(self.before), str(self.after), str(self.undoNext)]
```

```
1  /* Class for pages
2  class Page:
3      def __init__(self, pageId, data):
4          self.pageId = pageId
5          self.pageLSN = None
6          self.recLSN = None
7          self.data = data
8
9      def get_page(self):
10         return [str(self.pageId), str(self.pageLSN), str(self.recLSN), str(self.data)]
11
12 /* Class for transactions in Active Transaction Table (ATT)
13 class ActiveTransaction:
14     def __init__(self, id, status, lastLSN):
15         self.transactionId = id
16         self.status = status
17         self.lastLSN = lastLSN
18
19     def get_transaction(self):
20         return [str(self.transactionId), str(self.status), str(self.lastLSN)]
```



```
1  /* Represents the data on DISK(Non-volatile)
2  DISK_MEMORY = {
3      '1': Page('1', {'A': 5, 'B': 15}),
4      '2': Page('2', {'C': 10, 'D': 6}),
5      '3': Page('3', {'E': 20, 'F': 6}),
6  }
7
8  /* Represents the data on MEMORY(Volatile)
9  MEMORY = {}
10
11 /* Mapping variables to the pages they belong
12 page_mapping = {
13     'A': '1',
14     'B': '1',
15     'C': '2',
16     'D': '2',
17     'E': '3',
18     'F': '3'
19 }
20
21 /* Tempory buffer for modified data_variables
22 temp_buffer = {}
23
24 /* Active Transaction Table(ATT)
25 ATT = {}
26
27 /* Dirty Page Table (DPT)
28 DPT = {}
```

```

1  /* state = 1 (view before crash), state = 0 (view after recovery)
2  def view_system_state(state):
3      if (state == 1):
4          print('##### SYSTEM STATE BEFORE CRASH #####')
5
6      if (state == 0):
7          print('##### SYSTEM STATE AFTER RECOVERY #####')
8
9      main_records = []
10     print("  Log Records[On MEMORY]")
11     for LSN, logR in LOG_MEMORY.items():
12         main_records.append(logR.get_record())
13     print(tabulate(main_records, headers = ["LSN", "prevLSN", "TxnId", "type", "data", "before", "after", "undoNext"]))
14
15     print()
16     disk_records = []
17     print("  Log Records[On DISK]")
18     for LSN, logR in LOG_DISK.items():
19         disk_records.append(logR.get_record())
20     print(tabulate(disk_records, headers = ["LSN", "prevLSN", "TxnId", "type", "data", "before", "after", "undoNext"]))
21
22     print()
23     print("  MEMORY [Volatile]")
24     pages = []
25     for pageId, page in MEMORY.items():
26         pages.append(page.get_page())
27     print(tabulate(pages, ["pageId", "pageLSN", "recLSN", "data"]))
28
29     print()
30     print("  DISK MEMORY [NON-Volatile]")
31     pages = []
32     for pageId, page in DISK_MEMORY.items():
33         pages.append(page.get_page())
34     print(tabulate(pages, ["pageId", "pageLSN", "recLSN", "data"]))
35
36     print()
37     print("  Dirty Page Table(DPT)")
38     dirty_pages = []
39     for pageId, recLSN in DPT.items():
40         dirty_pages.append([pageId, recLSN])
41     print(tabulate(dirty_pages, headers = ["PageId", "recLSN"]))
42
43     print()
44     print("Active Transaction Table(ATT)")
45     transactions = []
46     for transactionId, transaction in ATT.items():
47         transactions.append(transaction.get_transaction())
48     print(tabulate(transactions, headers = ["TransactionId", "status", "lastLSN"]))
49
50     print('#####')

```



```
1  import ds
2  import rec
3  import copy
4
5  def get_value(ins, a, b):
6      if (ins.find('+') > -1): return a + b
7      if (ins.find('-') > -1): return a - b
8      if (ins.find('*') > -1): return a * b
9      if (ins.find('/') > -1): return a / b
```



```
1  if __name__ == '__main__':
2      schedule = open('schedule_input.txt', 'r')
3      data = []
4
5      for instruction in schedule:
6          data.append(instruction.strip())
7
8      /* Getting all instructions in the schedule */
9      instructions = data[0:]
10
11     execute_instructions(instructions)
```

Start Instructions Execution:

```
1 def execute_instructions(instructions):
2     for ins in instructions:
3
4         /* Normal Operations
5         if (ins.find('(') != -1):
6             operation = ins[0:ins.find('(')]
7             if (operation == 'CRASH'):
8                 ds.view_system_state(1)
9                 ds.LOG_MEMORY.clear()
10                ds.DPT.clear()
11                ds.ATT.clear()
12                ds.MEMORY.clear()
13                rec.execute_recovery()
14
15            if (operation == 'CHECKPOINT'):
16
17                checkpointData = ds.CheckPoint(ds.ATT, ds.DPT)
18                ds.masterRecord[0] = ds.LSN
19                ds.masterRecord[1] = checkpointData
20
21                log = ds.LogRecord(ds.LSN, None, None, 'CHECKPOINT', '-', '-', '-', '-')
22                ds.LOG_MEMORY[ds.LSN] = log
23
24                ds.LSN += 1
25
26            /* Flushing pages from MEMORY TO DISK
27            # pages_to_remove = []
28            if ds.flushedLSN != None:
29                for pageId, recLSN in ds.DPT.items():
30                    if recLSN <= ds.flushedLSN:
31                        ds.DISK_MEMORY[pageId] = copy.deepcopy(ds.MEMORY[pageId])
32                        # ds.MEMORY[pageId].pageLSN = None
33                        # ds.MEMORY[pageId].recLSN = None
34                        # pages_to_remove.append(pageId)
35
36            /* Clearing out flushed pages from DPT
37            # for id in pages_to_remove:
38            #     ds.DPT.pop(id)
```

```

1  if (operation == "BEGIN"):
2      transactionId = ins[ins.find('(') + 1:ins.find(')')]
3      log = ds.LogRecord(ds.LSN, None, transactionId, 'BEGIN', '-', '-', '-', '-')
4      ds.LOG_MEMORY[ds.LSN] = log
5      ds.ATT[transactionId] = ds.ActiveTransaction(transactionId, 'UNDO', ds.LSN)
6      ds.LSN += 1
7
8      if (operation == 'READ'):
9          data_variable = ins[ins.find('(') + 1:ins.find(',')]
10         transactionId = ins[ins.find(' ') + 1:ins.find(')')]
11
12         pageId = ds.page_mapping[data_variable]
13         if ds.MEMORY.get(pageId):                                     /* Page already present in MEMORY
14             continue
15         else:
16             ds.MEMORY[pageId] = copy.deepcopy(ds.DISK_MEMORY[pageId])
17
18     if (operation == 'WRITE'):
19         data_variable = ins[ins.find('(') + 1:ins.find(',')]
20         transactionId = ins[ins.find(' ') + 1:ins.find(')')]
21
22         prevLSN = ds.ATT[transactionId].lastLSN
23         ds.ATT[transactionId].lastLSN = ds.LSN
24
25         pageId = ds.page_mapping[data_variable]
26         before = ds.MEMORY[pageId].data.get(data_variable)
27         after = ds.temp_buffer[data_variable]
28
29         /* Writing Update Log record to MEMORY
30         log = ds.LogRecord(ds.LSN, prevLSN, transactionId, 'UPDATE', data_variable, before, after, '-')
31         ds.LOG_MEMORY[ds.LSN] = log
32
33         /* Updating the pages in MEMORY
34         ds.MEMORY[pageId].data[data_variable] = after
35         ds.MEMORY[pageId].pageLSN = ds.LSN
36
37         if ds.MEMORY[pageId].recLSN == None:
38             ds.MEMORY[pageId].recLSN = ds.LSN
39
40         /* Adding pages to dirty page table (DPT)
41         if ds.DPT.get(pageId) == None:
42             ds.DPT[pageId] = ds.LSN
43
44         ds.LSN += 1

```



```

1  if (operation == 'COMMIT'):
2      transactionId = ins[ins.find('(') + 1:ins.find(')')]
3
4      # print('COMMITTING TRANSACTION ' + str(transactionId))
5
6      ds.ATT[transactionId].status = 'COMMIT'
7
8      prevLSN = ds.ATT[transactionId].lastLSN
9      ds.ATT[transactionId].lastLSN = ds.LSN
10
11     /* Writing COMMIT Log record to MEMORY */
12     log = ds.LogRecord(ds.LSN, prevLSN, transactionId, 'COMMIT', '-', '-', '-', '-')
13     ds.LOG_MEMORY[ds.LSN] = log
14
15     /* Flushing all logs from MEMORY to DISK */
16     for LSN, log in ds.LOG_MEMORY.items():
17         ds.LOG_DISK[LSN] = log
18     ds.LOG_MEMORY.clear()
19
20     ds.flushedLSN = ds.LSN
21     ds.LSN += 1
22
23
24  if (operation == 'ABORT'):
25      transactionId = ins[ins.find('(') + 1:ins.find(')')]
26
27      ds.ATT[transactionId].status = 'ABORT'
28
29      prevLSN = ds.ATT[transactionId].lastLSN
30      ds.ATT[transactionId].lastLSN = ds.LSN
31
32     /* Writing ABORT Log record to MEMORY */
33     log = ds.LogRecord(ds.LSN, prevLSN, transactionId, 'ABORT', '-', '-', '-', '-')
34     ds.LOG_MEMORY[ds.LSN] = log
35
36     ds.LSN += 1
37
38     prev = prevLSN
39     while prev != None:
40         log = ds.LOG_MEMORY[prev]
41         if log.type == 'UPDATE':
42             prevLSN = ds.ATT[transactionId].lastLSN
43             ds.ATT[transactionId].lastLSN = ds.LSN
44             clr_log = ds.LogRecord(ds.LSN, prevLSN, transactionId, 'CLR', log.dataItem, log.after, log.before, log.prevLSN)
45             ds.LOG_MEMORY[ds.LSN] = clr_log
46
47             pageId = ds.page_mapping[log.dataItem]
48             /* Updating the pages in MEMORY */
49             ds.MEMORY[pageId].data[log.dataItem] = log.before
50             ds.MEMORY[pageId].pageLSN = ds.LSN
51             ds.LSN += 1
52             prev = log.prevLSN
53
54     /* Flushing all logs from MEMORY to DISK */
55     for LSN, log in ds.LOG_MEMORY.items():
56         ds.LOG_DISK[LSN] = log
57     ds.LOG_MEMORY.clear()
58
59     ds.flushedLSN = ds.LSN
60     ds.LSN += 1

```



```
1  else:
2      updating_data = ins[0]
3      a = ins[ins.find('[') + 1]
4
5      /* Code to Refactor */
6      b = ''
7      index = ins.find(']') - 1
8      while (ins[index] != ' '):
9          b += ins[index]
10         index -= 1
11     b = b[::-1]
12     b = int(b)
13
14     pageId = ds.page_mapping[a]
15     a = ds.MEMORY[pageId].data.get(a)
16     updated_value = get_value(ins, a, b)
17     ds.temp_buffer[updating_data] = updated_value
```


Recovery Code (After System Crash): Analysis Phase:

```
1  import ds
2  import copy
3
4  def analysis_phase():
5      lsn = ds.masterRecord[0]
6      ds.ATT = copy.deepcopy(ds.masterRecord[1].ATT)
7      ds.DPT = copy.deepcopy(ds.masterRecord[1].DPT)
8
9      for pageId, recLSN in ds.DPT.items():
10         ds.MEMORY[pageId] = copy.deepcopy(ds.DISK_MEMORY[pageId])
11
12     while ds.LOG_DISK.get(lsn):
13         log = ds.LOG_DISK[lsn]
14
15         /* Add the transaction in ATT with status UNDO
16         if log.type == 'BEGIN':
17             ds.ATT[log.TxnId] = ds.ActiveTransaction(log.TxnId, 'UNDO', log.LSN)
18
19         /* Add the page in DPT if not present and set recLSN with current LSN
20         if log.type == 'UPDATE':
21             pageId = ds.page_mapping[log.dataItem]
22
23             if not ds.DPT.get(pageId):
24                 ds.DPT[pageId] = log.LSN
25
26             ds.MEMORY[pageId] = copy.deepcopy(ds.DISK_MEMORY[pageId])
27             ds.ATT[log.TxnId].lastLSN = log.LSN
28
29         if log.type == 'COMMIT':
30             ds.ATT[log.TxnId].status = 'COMMIT'
31             ds.ATT[log.TxnId].lastLSN = log.LSN
32
33         if log.type == 'ABORT':
34             ds.ATT[log.TxnId].status = 'ABORT'
35             ds.ATT[log.TxnId].lastLSN = log.LSN
36
37         if log.type == 'END':
38             ds.ATT.pop(log.TxnId)
39         lsn += 1
```

Redo Phase:

```
1 def redo_phase():
2     LSN = 100
3
4     /* Get Minimum LSN to start redo scanning
5     for pageId, recLSN in ds.DPT.items():
6         if LSN > recLSN: LSN = recLSN
7
8     /* Redo all the update instructions
9     while ds.LOG_DISK.get(LSN):
10         log = ds.LOG_DISK[LSN]
11         LSN += 1
12
13         if log.type == 'UPDATE' or log.type == 'CLR':
14             pageId = ds.page_mapping[log.dataItem]
15
16             /* Redo the records for pages whose recLSN < LSN
17             for pId, recLSN in ds.DPT.items():
18                 /* Checking for ignore conditions
19                 if pageId != pId or ds.DISK_MEMORY[pId].recLSN >= LSN:
20                     continue
21                 else:
22                     ds.MEMORY[pId].data[log.dataItem] = log.after
23                     ds.MEMORY[pId].pageLSN = log.LSN
24
25             /* Redo all records both UPDATE AND CLR
26             # ds.MEMORY[pageId].data[log.dataItem] = log.after
27             # ds.MEMORY[pageId].pageLSN = log.LSN
28
29     transactions_to_remove = []
30     for transactionId, transactionInfo in ds.ATT.items():
31         if transactionInfo.status == 'COMMIT' or transactionInfo.status == 'ABORT':
32             transactions_to_remove.append(transactionId)
33             prevLSN = ds.ATT[transactionId].lastLSN
34             log = ds.LogRecord(ds.LSN, prevLSN, transactionId, 'END', '-', '-', '-', '-')
35             ds.LOG_MEMORY[ds.LSN] = log
36             ds.LSN += 1
37
38     for tid in transactions_to_remove:
39         ds.ATT.pop(tid)
```

Undo Phase:

```
1 def undo_phase():
2
3     transactions_to_remove = []
4     for transactionId, transaction in ds.ATT.items():
5         LSN = transaction.lastLSN
6
7         while LSN != None:
8             log = ds.LOG_DISK[LSN]
9             if log.type == 'UPDATE':
10                 prevLSN = ds.ATT[transactionId].lastLSN
11                 ds.ATT[transactionId].lastLSN = ds.LSN
12
13                 /* Adding CLR's for undo Updates */
14                 clr_log = ds.LogRecord(ds.LSN, prevLSN, transactionId, 'CLR', log.dataItem,
log.after, log.before, log.prevLSN)
15                 ds.LOG_MEMORY[ds.LSN] = clr_log
16
17                 /* Undoing the operations on pages */
18                 pageId = ds.page_mapping[log.dataItem]
19                 ds.MEMORY[pageId].data[log.dataItem] = log.before
20                 ds.MEMORY[pageId].pageLSN = ds.LSN
21
22                 ds.LSN += 1
23                 LSN = log.prevLSN
24
25                 prevLSN = ds.ATT[transactionId].lastLSN
26                 ds.ATT[transactionId].lastLSN = ds.LSN
27                 log = ds.LogRecord(ds.LSN, prevLSN, transactionId, 'ABORT', '-', '-', '-', '-')
28
29                 ds.LOG_MEMORY[ds.LSN] = log
30
31                 /* Flushing all logs from MEMORY to DISK */
32                 for LSN, log in ds.LOG_MEMORY.items():
33                     ds.LOG_DISK[LSN] = log
34                 ds.LOG_MEMORY.clear()
35
36                 ds.flushedLSN = ds.LSN
37                 ds.LSN += 1
38
39                 /* Flushing pages from MEMORY TO DISK */
40                 pages_to_remove = []
41                 for pageId, recLSN in ds.DPT.items():
42                     if recLSN <= ds.flushedLSN:
43                         ds.DISK_MEMORY[pageId] = copy.deepcopy(ds.MEMORY[pageId])
44                         pages_to_remove.append(pageId)
45
46                 /* Clearing out flushed pages from DPT */
47                 for id in pages_to_remove:
48                     ds.DPT.pop(id)
49
50                 /* Writing the TXN-END log record */
51                 prevLSN = ds.ATT[transactionId].lastLSN
52                 log = ds.LogRecord(ds.LSN, prevLSN, transactionId, 'END', '-', '-', '-', '-')
53                 ds.LOG_MEMORY[ds.LSN] = log
54
55                 transactions_to_remove.append(transactionId)
56                 ds.LSN += 1
57
58         for tid in transactions_to_remove:
59             ds.ATT.pop(tid)
```



```
1 def execute_recovery():  
2     /* Start the analysis Phase To create ATT and DPT  
3     analysis_phase()  
4     redo_phase()  
5     undo_phase()  
6  
7     ds.view_system_state(0)
```